

akshaya-project-13

April 28, 2023

#

Google App store analysis

0.1 Google App Store Data Analysis

Project Title : Google App Store Data Analysis

Technologies : Data Science

Domain : Technology

Project Difficulties level : Intermediate

0.2 Author : Akshaya L

```
[334]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
import re
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.feature_extraction.text import CountVectorizer
```

0.3 Basic analysis

First let's do some basic analysis to discover what our data looks like

```
[335]: df = pd.read_csv('./googleplaystore.csv')
df.head()
```

```
[335]:
```

	App	Category	Rating	\
0	Photo Editor & Candy Camera & Grid & ScrapBook	ART_AND_DESIGN	4.1	
1	Coloring book moana	ART_AND_DESIGN	3.9	

2	U Launcher Lite - FREE Live Cool Themes, Hide ...	ART_AND_DESIGN	4.7
3	Sketch - Draw & Paint	ART_AND_DESIGN	4.5
4	Pixel Draw - Number Art Coloring Book	ART_AND_DESIGN	4.3

	Reviews	Size	Installs	Type	Price	Content	Rating	\
0	159	19M	10,000+	Free	0		Everyone	
1	967	14M	500,000+	Free	0		Everyone	
2	87510	8.7M	5,000,000+	Free	0		Everyone	
3	215644	25M	50,000,000+	Free	0		Teen	
4	967	2.8M	100,000+	Free	0		Everyone	

	Genres	Last Updated	Current Ver	\
0	Art & Design	January 7, 2018	1.0.0	
1	Art & Design;Pretend Play	January 15, 2018	2.0.0	
2	Art & Design	August 1, 2018	1.2.4	
3	Art & Design	June 8, 2018	Varies with device	
4	Art & Design;Creativity	June 20, 2018	1.1	

	Android Ver
0	4.0.3 and up
1	4.0.3 and up
2	4.0.3 and up
3	4.2 and up
4	4.4 and up

We can already notice that we have a lot of categorical columns: Category, Installs, Type, Content Rating, Genres, Current Ver and Android Ver.

We can also notice the App column that contains the App name, meaning that all its values will most likely be unique.

```
[336]: df.duplicated().sum()
```

```
[336]: 483
```

We see that we have duplicated rows, let's drop them to not impact our analysis

```
[337]: df.shape
```

```
[337]: (10841, 13)
```

```
[338]: df.drop_duplicates(inplace=True)
df.shape
```

```
[338]: (10358, 13)
```

```
[339]: df.describe()
```

```
[339]:
```

	Rating
count	8893.000000
mean	4.189542
std	0.545452
min	1.000000
25%	4.000000
50%	4.300000
75%	4.500000
max	19.000000

Rating has a max of 19, while the play store rates from 1 to 5. We have invalid data to remove here.

```
[340]: df.sort_values(by=['Rating'], ascending=False).head()
```

```
[340]:
```

	App	Category	Rating	Reviews \
10472	Life Made WI-Fi Touchscreen Photo Frame	1.9	19.0	3.0M
5139	Chenoweth AH	MEDICAL	5.0	1
6851	BV Mobile Apps	PRODUCTIVITY	5.0	3
6807	Jabbla BT	TOOLS	5.0	3
6816	BU Study	FAMILY	5.0	7

	Size	Installs	Type	Price	Content	Rating	Genres \
10472	1,000+	Free	0	Everyone	NaN	February 11, 2018	
5139	27M	100+	Free	0	Everyone	Medical	
6851	4.8M	100+	Free	0	Everyone	Productivity	
6807	55k	100+	Free	0	Everyone	Tools	
6816	5.6M	10+	Free	0	Everyone	Education	

	Last Updated	Current Ver	Android Ver
10472	1.0.19	4.0 and up	NaN
5139	April 3, 2017	300000.0.78	4.0.3 and up
6851	June 5, 2018	2.0	4.2 and up
6807	October 6, 2014	1.0	4.2 and up
6816	December 7, 2017	1.0	4.0.3 and up

```
[341]: df.drop(10472, inplace=True);
```

```
[342]: df.shape
```

```
[342]: (10357, 13)
```

```
[343]: df.dtypes
```

```
[343]:
```

App	object
Category	object
Rating	float64
Reviews	object

```

Size          object
Installs      object
Type          object
Price         object
Content Rating object
Genres        object
Last Updated  object
Current Ver   object
Android Ver   object
dtype: object

```

We notice here that even if some columns like Reviews contains numeric values, they are encoded as objects.

```
[344]: df.isnull().sum()
```

```

[344]: App          0
       Category     0
       Rating      1465
       Reviews      0
       Size         0
       Installs     0
       Type         1
       Price        0
       Content Rating 0
       Genres       0
       Last Updated  0
       Current Ver   8
       Android Ver   2
       dtype: int64

```

Our data is pretty clean in term of NAN values, except for the Rating column which has 14% missing values

0.4 Cleaning up the data from any NAN values

```

[345]: # The columns 'Type', 'Content Rating', 'Current Ver' and 'Android Ver' have so
       ↪ few missing numbers
       # that we can delete these rows
       df_no_nan = df.dropna(axis=0, subset=['Type', 'Content Rating', 'Current Ver',
       ↪ 'Android Ver'])

```


```

[346]: # In terms of Rating, it all depends on what information we want to extract
       ↪ from our data. Let's see its distribution
       plt.figure(figsize=(5,5))
       bins = [0.5, 1.5, 2.5, 3.5, 4.5, 5.5, 6.5]
       df['Rating'].hist(bins = bins);

```

```
plt.xlim(0,7)
```

[346]: (0, 7)



We have a left skewed normal distribution with a maximum around 4.

We can already say that our data shows that people who DO rate are mostly the one that are the most satisfied with the apps. Thus the missing values might indicate an app that is not good enough for people to take the time to rate it. In this condition we cannot afford to take the decision to replace a missing value by the mean (which is 4.19) since it would probably not represent the reality.


We could then choose to either remove the missing value rows or give them a value like 0 which is outside of the rating range but would probably be more realistic of the type of grade we would get if the value was existing.

```
[347]: fill_nan = lambda col: col.fillna(0)
df_0_Rating = df_no_nan.apply(fill_nan, axis=0)

df_no_nan = df_no_nan.dropna(axis=0, subset=['Rating'])
```

0.5 Data analysis


```
[348]: plt.figure(figsize=(7,7))
plt.scatter(df_no_nan['Category'], df_no_nan['Rating'], alpha = .1);
plt.title('Ratings per category');
plt.xticks(rotation = 90);
```



We see here that the rating is mostly concentrated between 4 and 4.5, except for categories like family which have a larger range because they have more values.

```
[349]: df_no_nan['num_Installs'] = df_no_nan['Installs'].apply({ '10,000+':10000,
↳ '500,000+':500000, '5,000,000+':5000000, '50,000,000+':50000000, '100,000+':
↳ 100000, '50,000+':50000, '1,000,000+':1000000, '10,000,000+':10000000,
↳ '5,000+':5000, '100,000,000+':100000000, '1,000,000,000+':1000000000,
↳ '1,000+':1000, '500,000,000+':500000000, '50+':50, '100+':100, '500+':500,
↳ '10+':10, '1+':1, '5+':5, '0+':0 }.get)


sns.barplot(x='Category', y='num_Installs', data=df_no_nan, ci = None);
plt.title('Installs per category');
plt.xticks(rotation = 90);
```



We can see here that the communication apps are much more installed than any other categories


```
[350]: df_no_nan['Reviews'] = df_no_nan['Reviews'].astype('float64')

sns.barplot(x='Category', y='Reviews', data=df_no_nan, ci = None);
plt.title('Reviews per category');
plt.xticks(rotation = 90);
```



The social apps are the most reviewed, closely followed by the communication apps.

```
[351]: df_no_nan['Perc_Reviewed'] = df_no_nan['Reviews']/df_no_nan['num_Installs']*100;
sns.barplot(x='Category', y='Perc_Reviewed', data=df_no_nan.
    ↪sort_values(by='Perc_Reviewed', ascending=False), ci = None);
plt.title('Reviews per category (%)');
plt.xticks(rotation = 90);
```

The number of reviews per category is biased by the number of installations per category. If a category is more installed, it is expected to be more reviewed. To remove this bias we plotted the percentage of reviewed and now we see that the games are much likely to be reviewed than any other apps. Though they are closely followed by the social and medical apps.

```
[352]: df_number_per_category = df_no_nan.groupby('Category')['App'].nunique()
df_number_per_category = df_number_per_category.to_frame()
df_number_per_category['Category'] = df_number_per_category.index
df_number_per_category.columns=['Count', 'Category']
df_number_per_category['Percentage'] = df_number_per_category['Count']/
↳df_number_per_category.shape[0]
df_number_per_category.sort_values(by='Percentage', ascending=False).head(10)
```


```
[352]:
```

Category	Count	Category	Percentage
FAMILY	1683	FAMILY	51.000000
GAME	913	GAME	27.666667
TOOLS	719	TOOLS	21.787879
FINANCE	302	FINANCE	9.151515

LIFESTYLE	301	LIFESTYLE	9.121212
PRODUCTIVITY	301	PRODUCTIVITY	9.121212
PERSONALIZATION	296	PERSONALIZATION	8.969697
MEDICAL	291	MEDICAL	8.818182
BUSINESS	263	BUSINESS	7.969697
PHOTOGRAPHY	263	PHOTOGRAPHY	7.969697


Here we can see that 51% of the apps in the app store are classified as family while 27% are games and 21% are tools. Let's not forget that a lot of apps are classified in several categories and so the total percentage WILL BE higher than 100%

```
[353]: sns.barplot(x='Category', y='Percentage', data=df_number_per_category.
    ↪sort_values(by='Percentage', ascending=False), ci = None);
plt.title('App count per category (% of total)');
plt.xticks(rotation = 90);
```




```
[354]: df_no_nan['Price'] = df_no_nan['Price'].apply(lambda x : x.replace('$', ''))
df_no_nan['Price'] = df_no_nan['Price'].astype('float64')
plt.figure(figsize=(30, 10))
```

```
plt.title('Number of installation per price');
plt.xticks(rotation = 90, fontsize=14);
plt.yticks(fontsize=14);
sns.barplot(x='Price', y='num_Installs', data=df_no_nan, ci = None);
```



Here we can clearly see that the free apps are by far the most installed. After that, the apps paid 6.99\$ are the most installed

```
[355]: plt.figure(figsize=(30, 10))
plt.title('Number of installation per price');
plt.xticks(rotation = 90, fontsize=14);
plt.yticks(fontsize=14);
plt.yscale('log')
sns.barplot(x='Price', y='num_Installs', data=df_no_nan, ci = None);
```




```
[356]: df_no_nan['cumul_Installs_per_price'] = df_no_nan.
        ↳groupby(['Price'])['num_Installs'].apply(lambda x: x.cumsum())
```

```
df_no_nan_cumsum = df_no_nan.groupby(['Price']).agg({'cumul_Installs_per_price':
    ↪ 'sum'})
df_no_nan_cumsum['percentage']=df_no_nan_cumsum['cumul_Installs_per_price']/
    ↪df_no_nan_cumsum['cumul_Installs_per_price'].sum()*100
df_no_nan_cumsum.sort_values(by='percentage', ascending = False).head()
```

```
[356]:      cumul_Installs_per_price  percentage
Price
0.00      835889058796914      99.999711
0.99      1407569111      0.000168
2.99      481841680      0.000058
6.99      185293300      0.000022
4.99      102784560      0.000012
```

More than 99.99% of the installed apps are free!


```
[357]: plt.figure(figsize=(30, 10))
plt.title('Price per category');
plt.xticks(rotation = 90, fontsize=14);
plt.yticks(fontsize=14);
sns.barplot(x='Category', y='Price', data=df_no_nan, ci = None);
```



The price of the apps clearly depends on the category: the financial apps are sold at a much higher price than most of the other ones.


```
[358]: plt.figure(figsize=(30, 10))
plt.title('Count per genre');
plt.xticks(rotation = 90, fontsize=12);
plt.yticks(fontsize=12)
```

```
sns.barplot(y=df['Genres'].value_counts().reset_index()['Genres'],  
            x=df['Genres'].value_counts().reset_index()[:]['index']);
```



The apps with a genre tools, entertainment and education represent the most part of the play store.

```
[359]: plt.figure(figsize=(5, 5))
plt.title('Count per category');
plt.xticks(rotation = 90, fontsize=9);
plt.yticks(fontsize=9);
sns.barplot(y=df['Category'].value_counts().reset_index()['Category'],
            x=df['Category'].value_counts().reset_index()['index']);
```



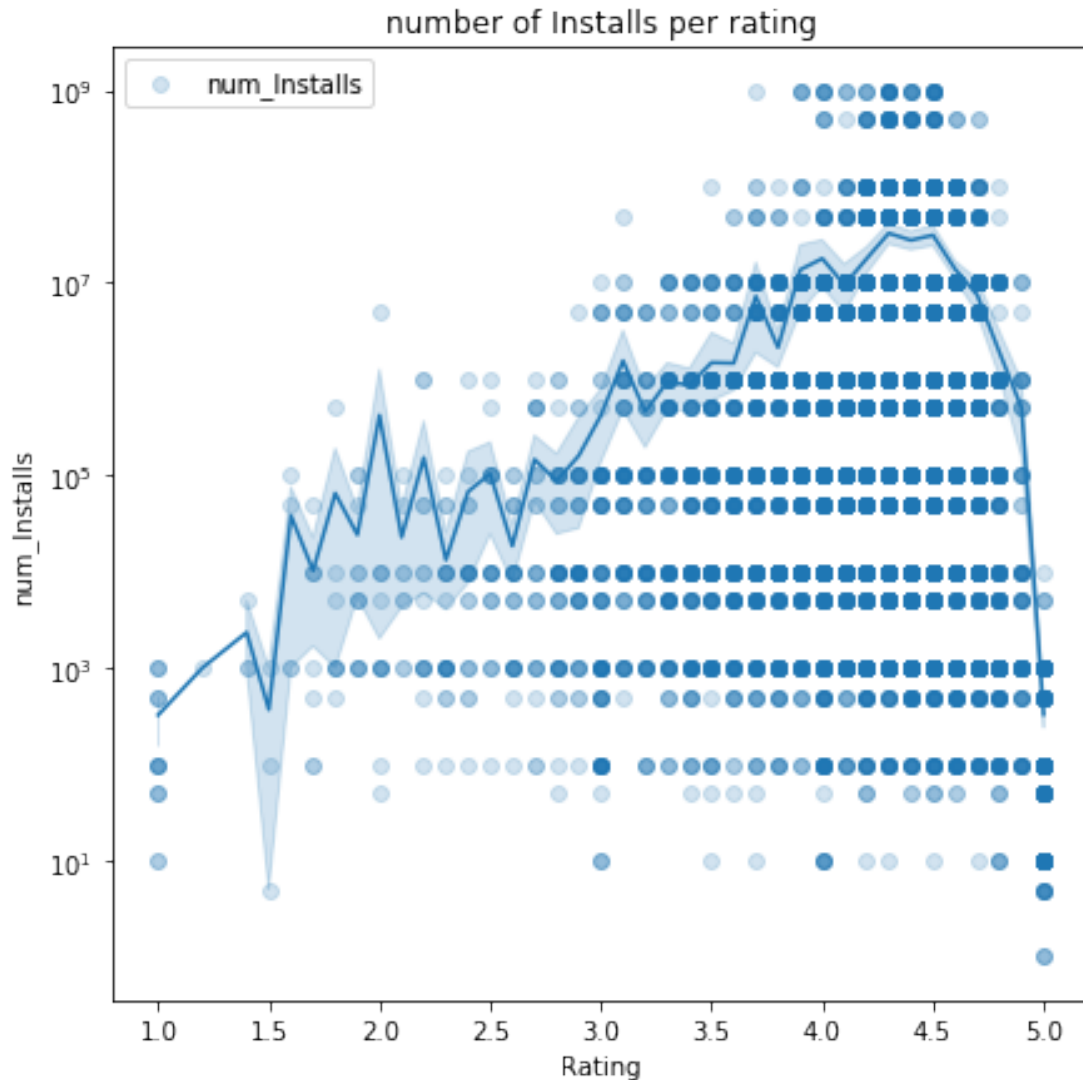
The categories the most represented is the app store are family, game and tools

```
[360]: plt.figure(figsize=(7,7))
plt.title('number of Installs per rating')
plt.scatter( x=df_no_nan['Rating'], y=df_no_nan['num_Installs'], alpha = 0.2)
sns.lineplot(x="Rating", y="num_Installs", data=df_no_nan)
plt.yscale('log')
```

C:\Users\Soizic\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713:
FutureWarning: Using a non-tuple sequence for multidimensional indexing is
deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will
be interpreted as an array index, `arr[np.array(seq)]`, which will result either


in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




We can see here that the number of installs and the rating are proportional until a rating of approximately 4.5, After that the number of installs drops considerably, probably meaning that the apps with a rating greater than 4.5 are mostly installed by friends and family

```
[361]: sns.barplot(x='num_Installs', y='Rating', data=df_no_nan, ci = None);  
plt.title('Rating per installs');  
plt.xticks(rotation = 90);
```




```
[362]: sns.barplot(x='Content Rating', y='Rating', data=df_no_nan, ci = None);  
plt.title('Rating per content rating');  
plt.xticks(rotation = 90);
```

The content rating of an app does not impact much its rating in the play store

```
[363]: plt.figure(figsize=(5,5))
plt.title('Number of installs per content rating')
sns.barplot(y="num_Installs", x="Content Rating", data=df_no_nan, ci=None)
plt.xticks(rotation = 90);
```



Though the apps targetting teenagers are much more installed than the other ones.

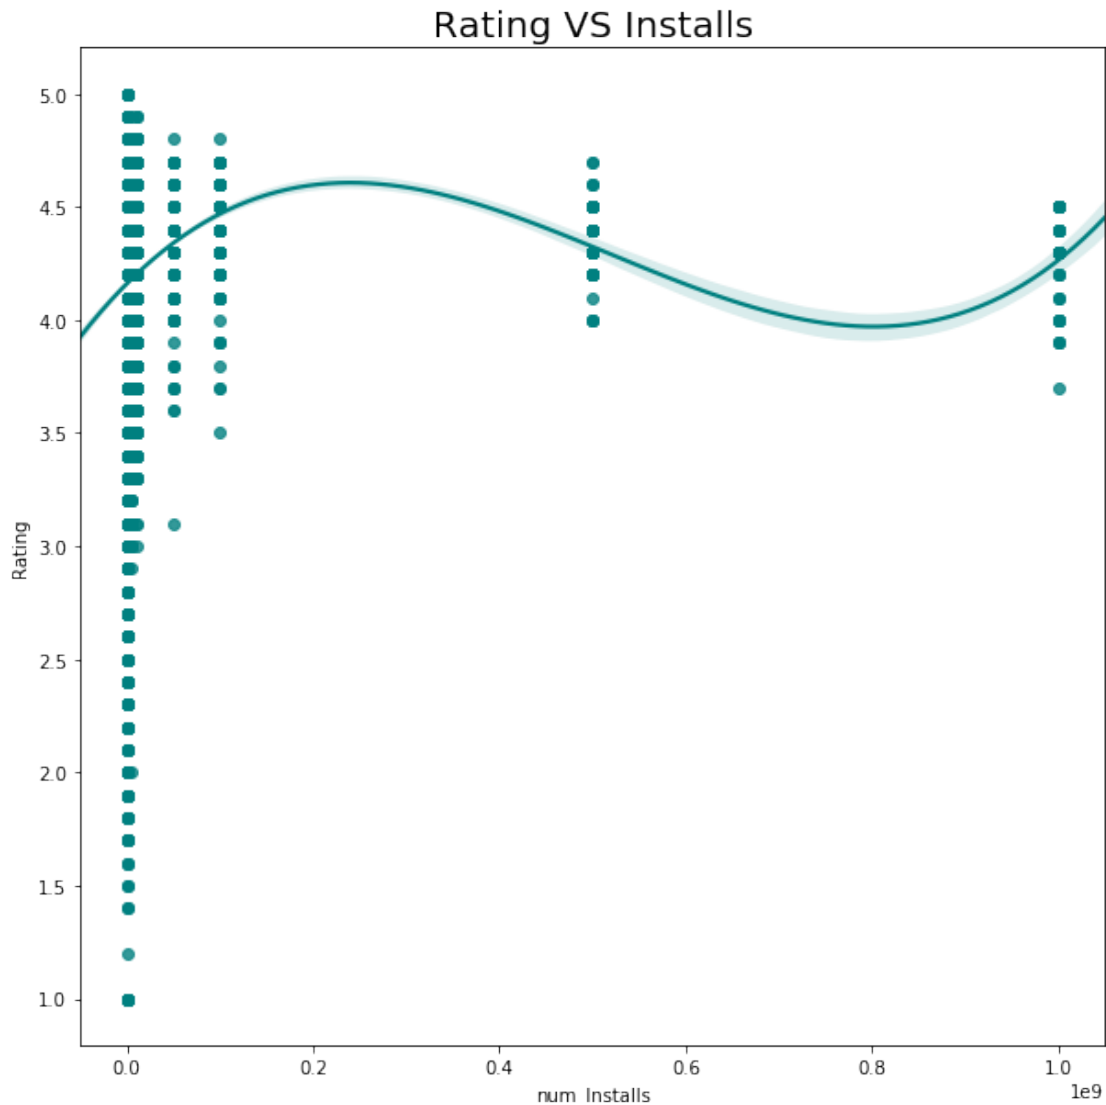
0.6 Modeling

```
[364]: plt.figure(figsize = (10,10))
sns.regplot(x="num_Installs", y="Rating", color = 'teal',data=df_no_nan,
            order=3);
plt.title('Rating VS Installs',size = 20)
```

C:\Users\Soizic\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713:
FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
[364]: Text(0.5,1,'Rating VS Installs')
```




```
[365]: def r2(x, y):  
        return stats.pearsonr(x, y)[0] ** 2  
sns.jointplot(df_no_nan['num_Installs'], df_no_nan['Rating'], kind="reg",  
             stat_func=r2, order=3)
```

```
C:\Users\Soizic\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713:  
FutureWarning: Using a non-tuple sequence for multidimensional indexing is  
deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will  
be interpreted as an array index, `arr[np.array(seq)]`, which will result either  
in an error or a different result.  
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
C:\Users\Soizic\Anaconda3\lib\site-packages\seaborn\axisgrid.py:1847:
UserWarning: JointGrid annotation is deprecated and will be removed in a future
release.
```

```
warnings.warn(UserWarning(msg))
```

```
[365]: <seaborn.axisgrid.JointGrid at 0x1842df9e0b8>
```



With an r^2 close to 0, we can definitely conclude that trying to approximate the correlation between ratings and the number of installs with a polynomial regression of rank 3 is not the right approach

```
[ ]:
```

```
[366]: df_no_nan.head()
```

```
[366]:
```

	App	Category	Rating \
0	Photo Editor & Candy Camera & Grid & ScrapBook	ART_AND_DESIGN	4.1
1	Coloring book moana	ART_AND_DESIGN	3.9
2	U Launcher Lite - FREE Live Cool Themes, Hide ...	ART_AND_DESIGN	4.7
3	Sketch - Draw & Paint	ART_AND_DESIGN	4.5
4	Pixel Draw - Number Art Coloring Book	ART_AND_DESIGN	4.3

	Reviews	Size	Installs	Type	Price	Content Rating \
0	159.0	19M	10,000+	Free	0.0	Everyone
1	967.0	14M	500,000+	Free	0.0	Everyone
2	87510.0	8.7M	5,000,000+	Free	0.0	Everyone
3	215644.0	25M	50,000,000+	Free	0.0	Teen
4	967.0	2.8M	100,000+	Free	0.0	Everyone

	Genres	Last Updated	Current Ver \
0	Art & Design	January 7, 2018	1.0.0
1	Art & Design;Pretend Play	January 15, 2018	2.0.0
2	Art & Design	August 1, 2018	1.2.4
3	Art & Design	June 8, 2018	Varies with device
4	Art & Design;Creativity	June 20, 2018	1.1

	Android Ver	num_Installs	Perc_Reviewed	cumul_Installs_per_price
0	4.0.3 and up	10000	1.590000	10000
1	4.0.3 and up	500000	0.193400	510000
2	4.0.3 and up	5000000	1.750200	5510000
3	4.2 and up	50000000	0.431288	55510000
4	4.4 and up	100000	0.967000	55610000

```
[367]: df_no_nan.drop(['App', 'Installs', 'Type', 'Last Updated', 'Current Ver', 'Android_Ver'], axis=1, inplace = True)
```

```
[368]: df_no_nan['Size'].unique()
```

```
[368]: array(['19M', '14M', '8.7M', '25M', '2.8M', '5.6M', '29M', '33M', '3.1M',
        '28M', '12M', '20M', '21M', '37M', '5.5M', '17M', '39M', '31M',
        '4.2M', '23M', '6.0M', '6.1M', '4.6M', '9.2M', '5.2M', '11M',
        '24M', 'Varies with device', '9.4M', '15M', '10M', '1.2M', '26M',
        '8.0M', '7.9M', '56M', '57M', '35M', '54M', '201k', '3.6M', '5.7M',
        '8.6M', '2.4M', '27M', '2.7M', '2.5M', '7.0M', '16M', '3.4M',
        '8.9M', '3.9M', '2.9M', '38M', '32M', '5.4M', '18M', '1.1M',
        '2.2M', '4.5M', '9.8M', '52M', '9.0M', '6.7M', '30M', '2.6M',
        '7.1M', '22M', '6.4M', '3.2M', '8.2M', '4.9M', '9.5M', '5.0M',
        '5.9M', '13M', '73M', '6.8M', '3.5M', '4.0M', '2.3M', '2.1M',
        '42M', '9.1M', '55M', '23k', '7.3M', '6.5M', '1.5M', '7.5M', '51M',
        '41M', '48M', '8.5M', '46M', '8.3M', '4.3M', '4.7M', '3.3M', '40M',
        '7.8M', '8.8M', '6.6M', '5.1M', '61M', '66M', '79k', '8.4M',
        '3.7M', '118k', '44M', '695k', '1.6M', '6.2M', '53M', '1.4M',
```

```

'3.0M', '7.2M', '5.8M', '3.8M', '9.6M', '45M', '63M', '49M', '77M',
'4.4M', '70M', '9.3M', '8.1M', '36M', '6.9M', '7.4M', '84M', '97M',
'2.0M', '1.9M', '1.8M', '5.3M', '47M', '556k', '526k', '76M',
'7.6M', '59M', '9.7M', '78M', '72M', '43M', '7.7M', '6.3M', '334k',
'93M', '65M', '79M', '100M', '58M', '50M', '68M', '64M', '34M',
'67M', '60M', '94M', '9.9M', '232k', '99M', '624k', '95M', '8.5k',
'41k', '292k', '80M', '1.7M', '10.0M', '74M', '62M', '69M', '75M',
'98M', '85M', '82M', '96M', '87M', '71M', '86M', '91M', '81M',
'92M', '83M', '88M', '704k', '862k', '899k', '378k', '4.8M',
'266k', '375k', '1.3M', '975k', '980k', '4.1M', '89M', '696k',
'544k', '525k', '920k', '779k', '853k', '720k', '713k', '772k',
'318k', '58k', '241k', '196k', '857k', '51k', '953k', '865k',
'251k', '930k', '540k', '313k', '746k', '203k', '26k', '314k',
'239k', '371k', '220k', '730k', '756k', '91k', '293k', '17k',
'74k', '14k', '317k', '78k', '924k', '818k', '81k', '939k', '169k',
'45k', '965k', '90M', '545k', '61k', '283k', '655k', '714k', '93k',
'872k', '121k', '322k', '976k', '206k', '954k', '444k', '717k',
'210k', '609k', '308k', '306k', '175k', '350k', '383k', '454k',
'1.0M', '70k', '812k', '442k', '842k', '417k', '412k', '459k',
'478k', '335k', '782k', '721k', '430k', '429k', '192k', '460k',
'728k', '496k', '816k', '414k', '506k', '887k', '613k', '778k',
'683k', '592k', '186k', '840k', '647k', '373k', '437k', '598k',
'716k', '585k', '982k', '219k', '55k', '323k', '691k', '511k',
'951k', '963k', '25k', '554k', '351k', '27k', '82k', '208k',
'551k', '29k', '103k', '116k', '153k', '209k', '499k', '173k',
'597k', '809k', '122k', '411k', '400k', '801k', '787k', '50k',
'643k', '986k', '516k', '837k', '780k', '20k', '498k', '600k',
'656k', '221k', '228k', '176k', '34k', '259k', '164k', '458k',
'629k', '28k', '288k', '775k', '785k', '636k', '916k', '994k',
'309k', '485k', '914k', '903k', '608k', '500k', '54k', '562k',
'847k', '948k', '811k', '270k', '48k', '523k', '784k', '280k',
'24k', '892k', '154k', '18k', '33k', '860k', '364k', '387k',
'626k', '161k', '879k', '39k', '170k', '141k', '160k', '144k',
'143k', '190k', '376k', '193k', '473k', '246k', '73k', '253k',
'957k', '420k', '72k', '404k', '470k', '226k', '240k', '89k',
'234k', '257k', '861k', '467k', '676k', '552k', '582k', '619k'],
dtype=object)

```

```

[369]: df_no_nan['Size'] = df_no_nan['Size'].apply(lambda x : x.replace('k', '000'))
df_no_nan['Size'] = df_no_nan['Size'].apply(lambda x : x.replace('M', '000000'))
df_no_nan['Size'] = df_no_nan['Size'].apply(lambda x : re.sub(r"\.", "", x))
df_no_nan['Size'] = df_no_nan['Size'].apply(lambda x : x.replace('Varies with_
device', 'NAN'))
df_no_nan['Size'] = df_no_nan['Size'].astype('float64')
df_no_nan = df_no_nan.dropna(subset=['Size'])
df_no_nan['Size'].unique()

```

```
[369]: array([1.90e+07, 1.40e+07, 8.00e+06, 2.50e+07, 2.00e+06, 5.00e+06,
2.90e+07, 3.30e+07, 3.00e+06, 2.80e+07, 1.20e+07, 2.00e+07,
2.10e+07, 3.70e+07, 1.70e+07, 3.90e+07, 3.10e+07, 4.00e+06,
2.30e+07, 6.00e+06, 9.00e+06, 1.10e+07, 2.40e+07, 1.50e+07,
1.00e+07, 1.00e+06, 2.60e+07, 7.00e+06, 5.60e+07, 5.70e+07,
3.50e+07, 5.40e+07, 2.01e+05, 2.70e+07, 1.60e+07, 3.80e+07,
3.20e+07, 1.80e+07, 5.20e+07, 3.00e+07, 2.20e+07, 1.30e+07,
7.30e+07, 4.20e+07, 5.50e+07, 2.30e+04, 5.10e+07, 4.10e+07,
4.80e+07, 4.60e+07, 4.00e+07, 6.10e+07, 6.60e+07, 7.90e+04,
1.18e+05, 4.40e+07, 6.95e+05, 5.30e+07, 4.50e+07, 6.30e+07,
4.90e+07, 7.70e+07, 7.00e+07, 3.60e+07, 8.40e+07, 9.70e+07,
4.70e+07, 5.56e+05, 5.26e+05, 7.60e+07, 5.90e+07, 7.80e+07,
7.20e+07, 4.30e+07, 3.34e+05, 9.30e+07, 6.50e+07, 7.90e+07,
1.00e+08, 5.80e+07, 5.00e+07, 6.80e+07, 6.40e+07, 3.40e+07,
6.70e+07, 6.00e+07, 9.40e+07, 2.32e+05, 9.90e+07, 6.24e+05,
9.50e+07, 8.00e+03, 4.10e+04, 2.92e+05, 8.00e+07, 7.40e+07,
6.20e+07, 6.90e+07, 7.50e+07, 9.80e+07, 8.50e+07, 8.20e+07,
9.60e+07, 8.70e+07, 7.10e+07, 8.60e+07, 9.10e+07, 8.10e+07,
9.20e+07, 8.30e+07, 8.80e+07, 7.04e+05, 8.62e+05, 8.99e+05,
3.78e+05, 2.66e+05, 3.75e+05, 9.75e+05, 9.80e+05, 8.90e+07,
6.96e+05, 5.44e+05, 5.25e+05, 9.20e+05, 7.79e+05, 8.53e+05,
7.20e+05, 7.13e+05, 7.72e+05, 3.18e+05, 5.80e+04, 2.41e+05,
1.96e+05, 8.57e+05, 5.10e+04, 9.53e+05, 8.65e+05, 2.51e+05,
9.30e+05, 5.40e+05, 3.13e+05, 7.46e+05, 2.03e+05, 2.60e+04,
3.14e+05, 2.39e+05, 3.71e+05, 2.20e+05, 7.30e+05, 7.56e+05,
9.10e+04, 2.93e+05, 1.70e+04, 7.40e+04, 1.40e+04, 3.17e+05,
7.80e+04, 9.24e+05, 8.18e+05, 8.10e+04, 9.39e+05, 1.69e+05,
4.50e+04, 9.65e+05, 9.00e+07, 5.45e+05, 6.10e+04, 2.83e+05,
6.55e+05, 7.14e+05, 9.30e+04, 8.72e+05, 1.21e+05, 3.22e+05,
9.76e+05, 2.06e+05, 9.54e+05, 4.44e+05, 7.17e+05, 2.10e+05,
6.09e+05, 3.08e+05, 3.06e+05, 1.75e+05, 3.50e+05, 3.83e+05,
4.54e+05, 7.00e+04, 8.12e+05, 4.42e+05, 8.42e+05, 4.17e+05,
4.12e+05, 4.59e+05, 4.78e+05, 3.35e+05, 7.82e+05, 7.21e+05,
4.30e+05, 4.29e+05, 1.92e+05, 4.60e+05, 7.28e+05, 4.96e+05,
8.16e+05, 4.14e+05, 5.06e+05, 8.87e+05, 6.13e+05, 7.78e+05,
6.83e+05, 5.92e+05, 1.86e+05, 8.40e+05, 6.47e+05, 3.73e+05,
4.37e+05, 5.98e+05, 7.16e+05, 5.85e+05, 9.82e+05, 2.19e+05,
5.50e+04, 3.23e+05, 6.91e+05, 5.11e+05, 9.51e+05, 9.63e+05,
2.50e+04, 5.54e+05, 3.51e+05, 2.70e+04, 8.20e+04, 2.08e+05,
5.51e+05, 2.90e+04, 1.03e+05, 1.16e+05, 1.53e+05, 2.09e+05,
4.99e+05, 1.73e+05, 5.97e+05, 8.09e+05, 1.22e+05, 4.11e+05,
4.00e+05, 8.01e+05, 7.87e+05, 5.00e+04, 6.43e+05, 9.86e+05,
5.16e+05, 8.37e+05, 7.80e+05, 2.00e+04, 4.98e+05, 6.00e+05,
6.56e+05, 2.21e+05, 2.28e+05, 1.76e+05, 3.40e+04, 2.59e+05,
1.64e+05, 4.58e+05, 6.29e+05, 2.80e+04, 2.88e+05, 7.75e+05,
7.85e+05, 6.36e+05, 9.16e+05, 9.94e+05, 3.09e+05, 4.85e+05,
9.14e+05, 9.03e+05, 6.08e+05, 5.00e+05, 5.40e+04, 5.62e+05,
```

```
8.47e+05, 9.48e+05, 8.11e+05, 2.70e+05, 4.80e+04, 5.23e+05,
7.84e+05, 2.80e+05, 2.40e+04, 8.92e+05, 1.54e+05, 1.80e+04,
3.30e+04, 8.60e+05, 3.64e+05, 3.87e+05, 6.26e+05, 1.61e+05,
8.79e+05, 3.90e+04, 1.70e+05, 1.41e+05, 1.60e+05, 1.44e+05,
1.43e+05, 1.90e+05, 3.76e+05, 1.93e+05, 4.73e+05, 2.46e+05,
7.30e+04, 2.53e+05, 9.57e+05, 4.20e+05, 7.20e+04, 4.04e+05,
4.70e+05, 2.26e+05, 2.40e+05, 8.90e+04, 2.34e+05, 2.57e+05,
8.61e+05, 4.67e+05, 6.76e+05, 5.52e+05, 5.82e+05, 6.19e+05])
```

```
[370]: df_no_nan.select_dtypes(include=['object']).columns
```

```
[370]: Index(['Category', 'Content Rating', 'Genres'], dtype='object')
```

```
[371]: df_no_nan['Reviews'] = df_no_nan['Reviews'].astype('float64')
```

```
[372]: cat_vars = df_no_nan.select_dtypes(include=['object']).columns
for col in cat_vars:
    df_no_nan = pd.concat([df_no_nan.drop([col], axis=1), pd.
↳get_dummies(df_no_nan[col], prefix=col)], axis=1)
```

```
[373]: df_no_nan.head()
```

```
[373]:
```

	Rating	Reviews	Size	Price	num_Installs	Perc_Reviewed	\
0	4.1	159.0	19000000.0	0.0	10000	1.590000	
1	3.9	967.0	14000000.0	0.0	500000	0.193400	
2	4.7	87510.0	8000000.0	0.0	5000000	1.750200	
3	4.5	215644.0	25000000.0	0.0	50000000	0.431288	
4	4.3	967.0	2000000.0	0.0	100000	0.967000	

	cumul_Installs_per_price	Category_ART_AND_DESIGN	\
0	10000	1	
1	510000	1	
2	5510000	1	
3	55510000	1	
4	55610000	1	

	Category_AUTO_AND_VEHICLES	Category_BEAUTY	...	\
0	0	0	...	
1	0	0	...	
2	0	0	...	
3	0	0	...	
4	0	0	...	

	Genres_Strategy;Education	Genres_Tools	Genres_Travel & Local	\
0	0	0	0	
1	0	0	0	
2	0	0	0	

3	0	0	0
4	0	0	0

	Genres_Travel & Local;Action & Adventure	Genres_Trivia \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	Genres_Video Players & Editors	Genres_Video Players & Editors;Creativity \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	Genres_Video Players & Editors;Music & Video	Genres_Weather	Genres_Word
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 158 columns]

```
[374]: y = df_no_nan['Rating']
X = df_no_nan.drop(['Rating'], axis=1)
```

```
[375]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳random_state=42)

clf = RandomForestClassifier(n_estimators=100, max_depth=20,random_state=42)
lab_enc = preprocessing.LabelEncoder()
y_encoded = lab_enc.fit_transform(y_train)
clf.fit(X_train, y_encoded)

y_pred = clf.predict(X_test)
y_test_encoded = lab_enc.fit_transform(y_test)
accuracy_score(y_test_encoded, y_pred)
```

```
[375]: 0.10062893081761007
```


With an accuracy of 11% (the best I was able to achieve when playing with the parameters), we cannot really predict our rating based on the other information with have in the dataset.

```
[376]: importances = clf.feature_importances_  
indices = np.argsort(importances)[::-1]  
feature_names = df_no_nan.drop(['Rating'], axis=1).columns  
f, ax = plt.subplots(figsize=(20,20));  
plt.title("Feature ranking", fontsize = 20);  
plt.bar(range(X_train.shape[1]), importances[indices],align="center");  
plt.xticks(range(X_train.shape[1]), indices);  
plt.xlim([-1, X_train.shape[1]]);  
plt.ylabel("importance", fontsize = 18);  
plt.xlabel("index of the feature", fontsize = 18);  
plt.xticks(range(X_train.shape[1]), feature_names, rotation=90);
```



```
[377]: model = CountVectorizer(max_features=10, stop_words='english')
X = model.fit_transform(list(df['App']))

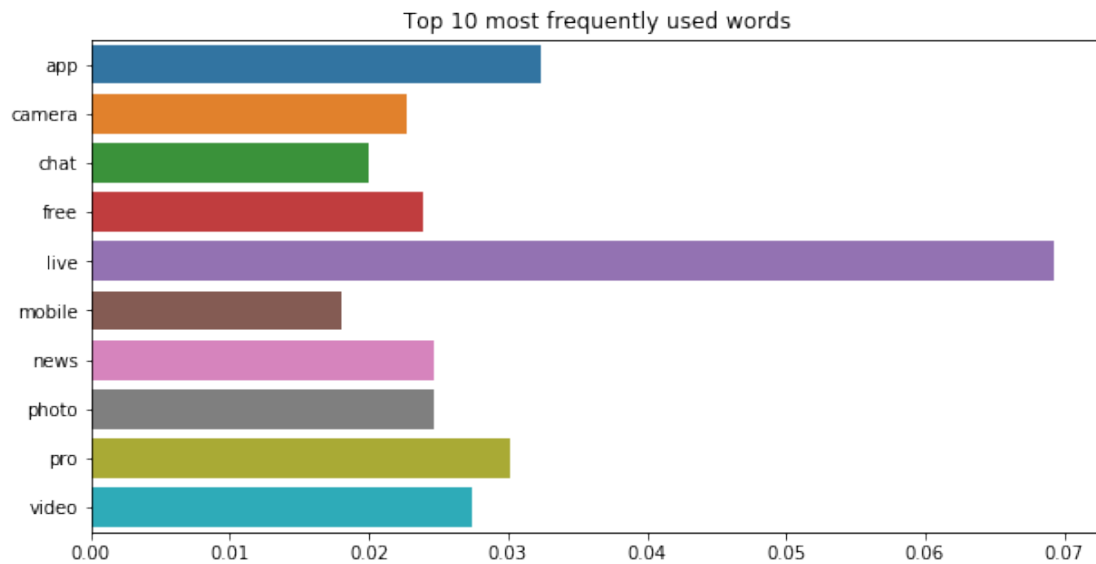
plt.figure(figsize=(10,5))
plt.title('Top 10 most frequently used words')
sns.barplot(x=X.toarray().mean(axis=0), y=vectorizer.get_feature_names());
```



More than 4% of the apps in the Google play store have the word free in their name! The 2 other most important keywords are app and mobile.

```
[378]: df['num_Installs'] = df['Installs'].apply({ '10,000+':10000, '500,000+':500000,
↪ '5,000,000+':5000000, '50,000,000+':50000000, '100,000+':100000, '50,000+':
↪ 50000, '1,000,000+':1000000, '10,000,000+':10000000, '5,000+':5000,
↪ '100,000,000+':100000000, '1,000,000,000+':1000000000, '1,000+':1000,
↪ '500,000,000+':500000000, '50+':50, '100+':100, '500+':500, '10+':10, '1+':
↪ 1, '5+':5, '0+':0 }.get)
X = model.fit_transform(list(df[df['num_Installs']>=1000000]['App']))

plt.figure(figsize=(10,5))
plt.title('Top 10 most frequently used words')
sns.barplot(x=X.toarray().mean(axis=0), y=vectorizer.get_feature_names());
```



But if we look at only at the apps that have been installed 1,000,000+ times, then the most important keyword is “live”.

[]: